


# Hidden Markov Models

Toby Dylan Hocking

# Background: detecting abrupt changes is important

Example from cancer diagnosis: breakpoints are associated with aggressive disease in neuroblastoma.




# Motivation for Hidden Markov Models (HMMs)

- ▶ Sometimes we have an interpretation / expectation of what the segments/clusters mean.
- ▶ For example in DNA copy number data the  $\text{logratio}=0$  means normal copy number (two copies – one from each parent), whereas higher  $\text{logratio}$  values indicate gain/amplification and lower values indicate loss/deletion.


# HMM ideas

- ▶ Each observed data variable in the sequence has a corresponding un-observed (hidden) state variable.
- ▶ There are typically a finite number of possible values for each hidden state variable,  $k \in \{1, \dots, K\}$ .
- ▶ Markov assumption: first-order dependency (each hidden variable only depends on the previous hidden variable in the sequence).



# Parameters of HMM

- ▶ transition matrix:  $A \in [0, 1]^{K \times K}$  for  $K$  clusters. Each entry  $a_{ij}$  is the probability of transitioning from state  $i$  to state  $j$ .
- ▶ initial state distribution:  $\pi \in [0, 1]^K$  (prior weights).
- ▶ emission: define the probability of observing data  $y$  in state  $k$ ,  $b_j(y) \in [0, 1]$  (mean, covariance).
- ▶ These parameters are unknown in advance and must be learned from the data.
- ▶ Comparison with Gaussian Mixture Models: HMM has all of GMM parameters, plus transition matrix.




# Three problems and algorithms

- ▶ Evaluation of likelihood of a given data sequence and model parameters  $\Rightarrow$  forward algorithm.
- ▶ Decoding most likely sequence of hidden states given observed data  $\Rightarrow$  Viterbi algorithm.
- ▶ Learning model parameters that maximize likelihood of a given data set  $\Rightarrow$  Baum-Welch algorithm.

# Forward algorithm

- ▶ Assume each state  $k \in \{1, \dots, K\}$  has an emission probability function  $b_k(o_t)$  for observed data  $o_t$  at time  $t$ .
- ▶ Let  $a_{kj}$  be a transition parameter (probability of going from state  $k$  to state  $j$ ).
- ▶ Then we can recursively compute the forward path probability,  $\alpha_t(j) = \sum_{k=1}^K \alpha_{t-1}(k) a_{kj} b_j(o_t)$ .
- ▶ Total log likelihood is summed over all states at the last data point in the sequence.



# Viterbi algorithm

- ▶ Assume each state  $k \in \{1, \dots, K\}$  has an emission probability function  $b_k(o_t)$  for observed data  $o_t$  at time  $t$ .
- ▶ Let  $a_{kj}$  be a transition parameter (probability of going from state  $k$  to state  $j$ ).
- ▶ Then we can recursively compute the probability of the best sequence of hidden variables that ends at data point  $t$  in state  $j$ ,  $v_t(j) = \max_{k \in \{1, \dots, K\}} v_{t-1}(k) a_{kj} b_j(o_t)$ .



# Viterbi algorithm example


```
## List of 4
```

```
## $ mean      : num [1:3] 0.25 0 -0.5
```

```
## $ sd        : num 1
```


```
## $ transition: num [1:3, 1:3] 0.333 0.333 0.333 0.333 0
```

```
## $ initial   : num [1:3] 0.333 0.333 0.333
```




# Baum-Welch learning algorithm


- ▶ Is an instance of Expectation-Maximization (EM), like the Gaussian Mixture Model learning algorithm.
- ▶ E step involves forward/backward passes over data sequences, to compute probability of each data point in each state.
- ▶ M step involves re-computing model parameters.
- ▶ Repeat until the log likelihood stops increasing.




# Visualization of learning iterations



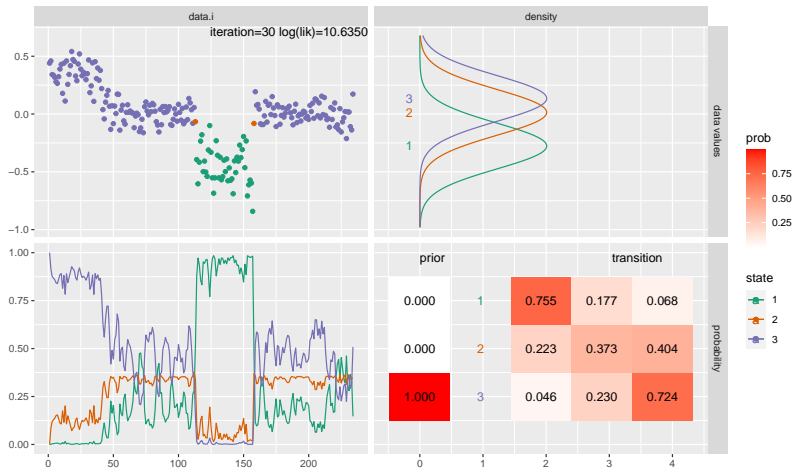
# Visualization of learning iterations




# Visualization of learning iterations




# Visualization of learning iterations



# Visualization of learning iterations




# Visualization of learning iterations






# Visualization of learning iterations




# Shared states between chromosomes on a profile?


- ▶ Sometimes there are several data sequences which are assumed to have the same set of hidden states.
- ▶ In this case Baum-Welch can be used to fit HMM to all data at the same time (total log likelihood is summed over all data sequences).




# HMM on whole profile




# Learning iterations using multiple sequences



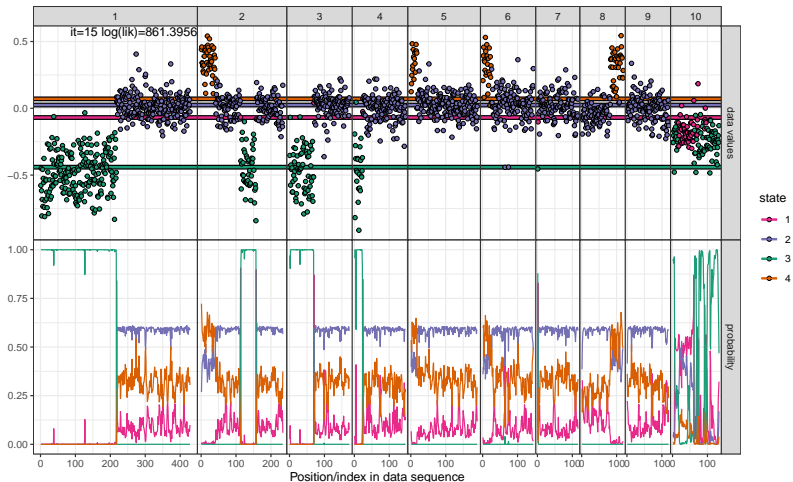
# Learning iterations using multiple sequences




# Learning iterations using multiple sequences



# Learning iterations using multiple sequences




# Learning iterations using multiple sequences






# Learning iterations using multiple sequences



# Learning iterations using multiple sequences



Time/space complexity TODO

## Possible exam questions TODO